

Introduction to Container & Docker

Agenda

- * **Container & Docker:**
- * **Introduction to Container & Docker:**
 - * Container Benefits.
 - * Docker Components.
 - * Lab Resources.
- * **Docker Engine:**
 - * Docker Basic Command.
 - * Labs: Docker Basic & Exercise.
- * **Docker File:**
 - * Usage Dockerfile.
 - * Dockerfile Syntax & Commands.
 - * Lab: Dockerfile & Exercise.
- * **Docker Compose:**
 - * Docker Compose Commands.
 - * Lab: Docker Compose & Exercise.
- * **Container orchestration platform**
- * **Automation and pipeline:**
 - * What's git?
 - * What's gitlab?
 - * Gitlab pipelines How can we implement Docker commands on gitlab pipelines?
 - * Best practices for gitlab (Branching and Merge requests).

Container

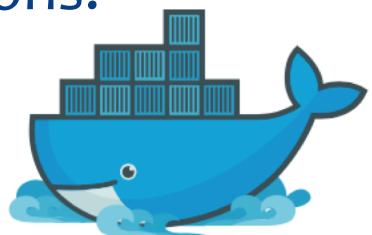
Containers are a lightweight and portable store for an **application** and its **dependencies**.

What's the Diff: VMs vs Containers

VMs	Containers
Heavyweight	Lightweight
Limited performance	Native performance
Each VM runs in its own OS	All containers share the host OS
Hardware-level virtualization	OS virtualization
Startup time in minutes	Startup time in milliseconds
Allocates required memory	Requires less memory space
Fully isolated and hence more secure	Process-level isolation, possibly less secure

Docker Inc

- * Founded in 2009.
- * Formerly dotCloud Inc.
- * Primary sponsor of the Docker Project.
- * Together with the community of maintainers and contributors, Docker aims to deliver open tools to help developers build applications with open APIs to help sysadmins better manage these applications.
- * Services :- Docker Cloud, Docker Datacenter

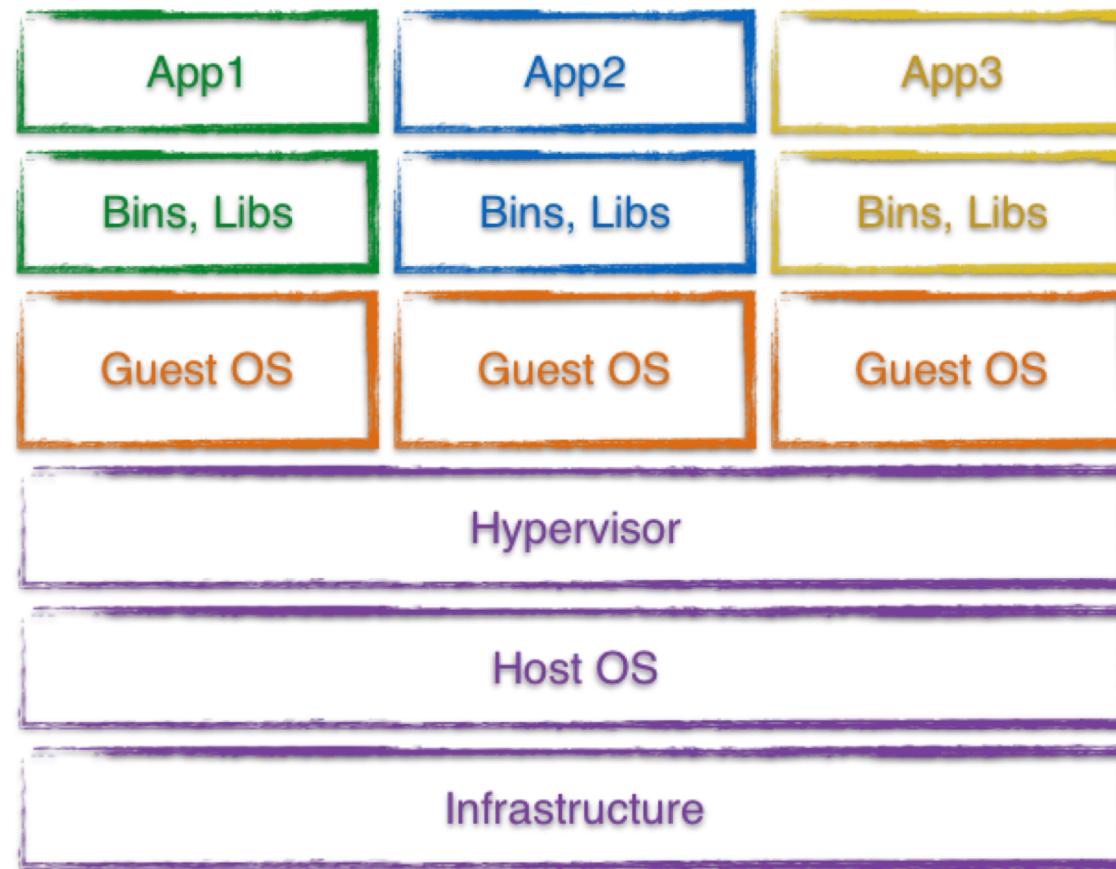


Container Benefits

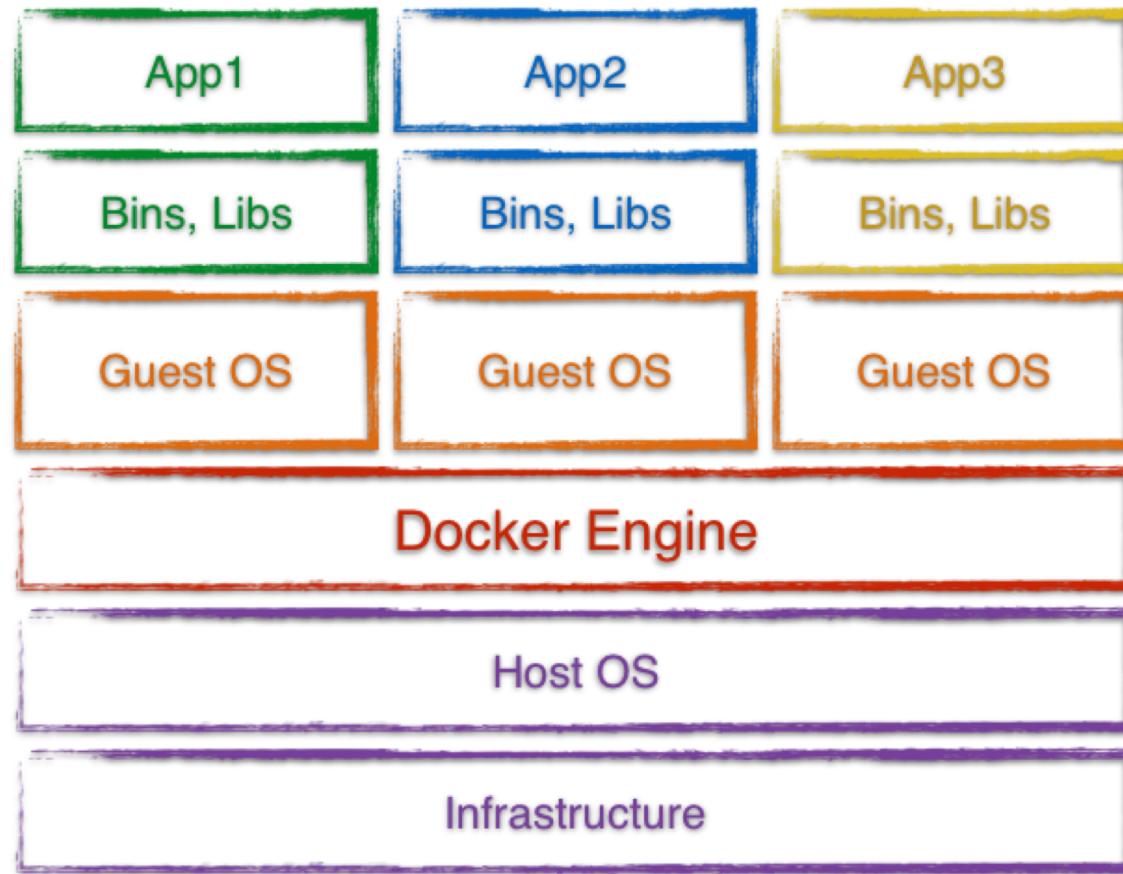
Containers are a streamlined way to build, test, deploy, and redeploy applications on multiple environments from a developer's local laptop to an on-premises data center and even the cloud. Benefits of containers include:

- * **Less overhead.** Containers require less system resources than traditional or hardware virtual machine environments because they don't include operating system images.
- * **Increased portability.** Applications running in containers can be deployed easily to multiple different operating systems and hardware platforms.
- * **More consistent operation.** DevOps teams know applications in containers will run the same, regardless of where they are deployed.
- * **Greater efficiency.** Containers allow applications to be more rapidly deployed, patched, or scaled.
- * **“But it works on my machine”** stuff should be gone.

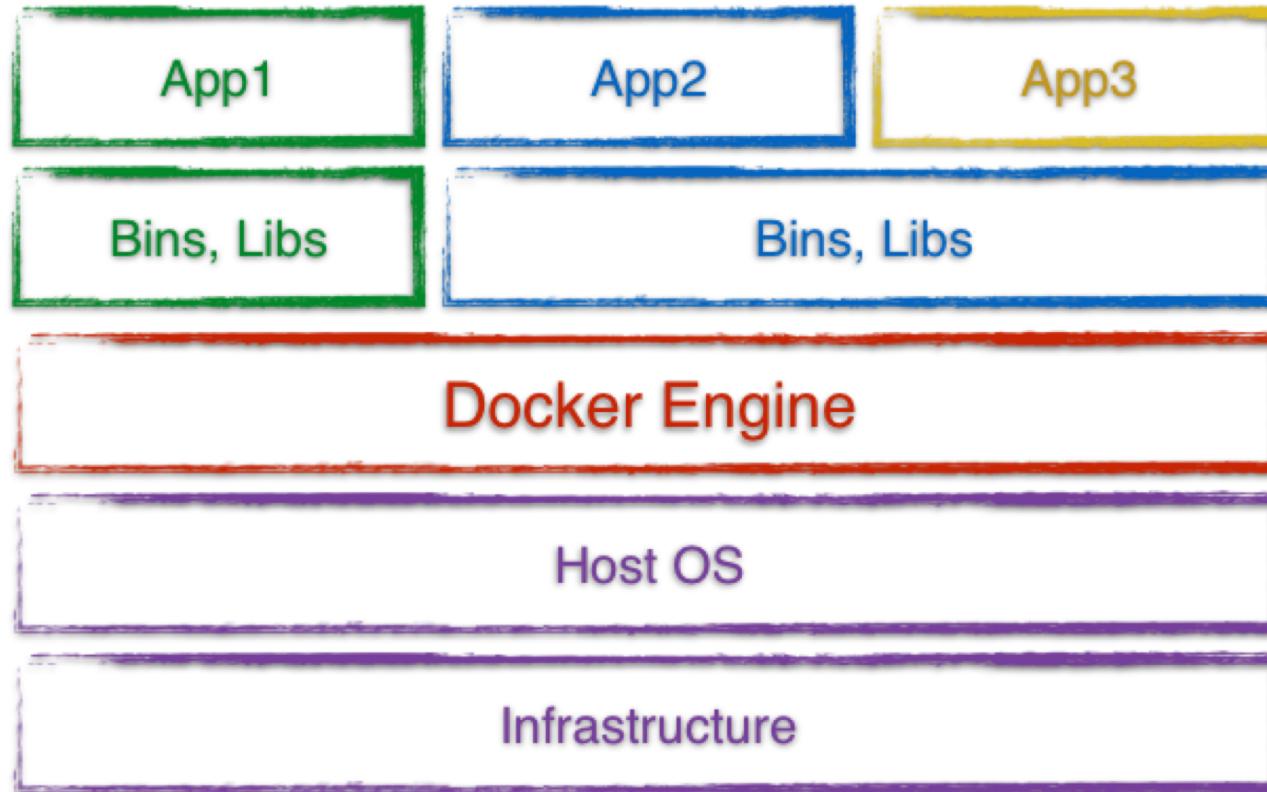
Virtual Machine



Container



Container



Docker Components



- * **Docker Engine :-**
Create Docker images and run Docker containers. Including **Swarm** mode from 1.12.0-rc1



- * **Docker Compose :-**
Defines applications built using multiple containers.



- * **Docker Hub :-**
A hosted registry service for managing and building images.



- * **Docker Cloud :-**
A hosted service for building, testing, and deploying Docker images to your hosts.



- * **Docker Trusted Registry :-**
(DTR) stores and signs your images.



- * **Docker Machine :-**
Automate container provisioning on your network or in the cloud. Available for Windows, Mac OS X, or Linux

Supported Platforms



Docker for Mac

A native application using the OS X sandbox security model which delivers all Docker tools to your Mac.



Docker for Windows

A native Windows application which delivers all Docker tools to your Windows computer.



Docker for Linux

Install Docker on a computer which already has a Linux distribution installed.

Your training Virtual Machine

#	IP	#	IP
1	10.254.145.226	11	10.254.145.218
2	10.254.145.141	12	10.254.145.175
3	10.254.145.10	13	10.254.145.88
4	10.254.145.20	14	10.254.145.214
5	10.254.145.223	15	10.254.145.232
6	10.254.145.231	16	10.254.145.212
7	10.254.145.96	17	10.254.145.180
8	10.254.145.140	18	10.254.145.225
9	10.254.145.71	19	10.254.145.224
10	10.254.145.78	20	10.254.145.203

Connecting to your Virtual Machine

- * You need an SSH client.
- * On OS X, Linux, and other UNIX systems, just use ssh:
`$ ssh <KAUST-ID>@<ip-address>`
- * On Windows, if you don't have an SSH client, you can download Putty from www.putty.org.

Checking your Virtual Machine

- * Once logged in, make sure that you can run a basic Docker command:

```
$ docker version
Client:
  Version:          18.09.5
  API version:      1.39
  Go version:       go1.10.8
  Git commit:       e8ff056
  Built:            Thu Apr 11 04:43:57 2019
  OS/Arch:          linux/amd64
  Experimental:    false
Server: Docker Engine - Community
Engine:
  Version:          18.09.5
  API version:      1.39 (minimum version 1.12)
  Go version:       go1.10.8
  Git commit:       e8ff056
  Built:            Thu Apr 11 04:10:53 2019
  OS/Arch:          linux/amd64
  Experimental:    false
```

Docker architecture

- * The Docker daemon (or "Engine")
Receives and processes incoming Docker API requests.
- * The Docker client
Talks to the Docker daemon via the Docker API. We'll
use mostly the CLI embedded within the docker binary.

Docker Basics

- * Check the version of docker

```
$ docker --version
```

- * Run the first docker container... Hello, World!

```
$ docker run hello-world
```

- * Check the process of docker container

```
$ docker ps
```

```
$ docker ps -a
```

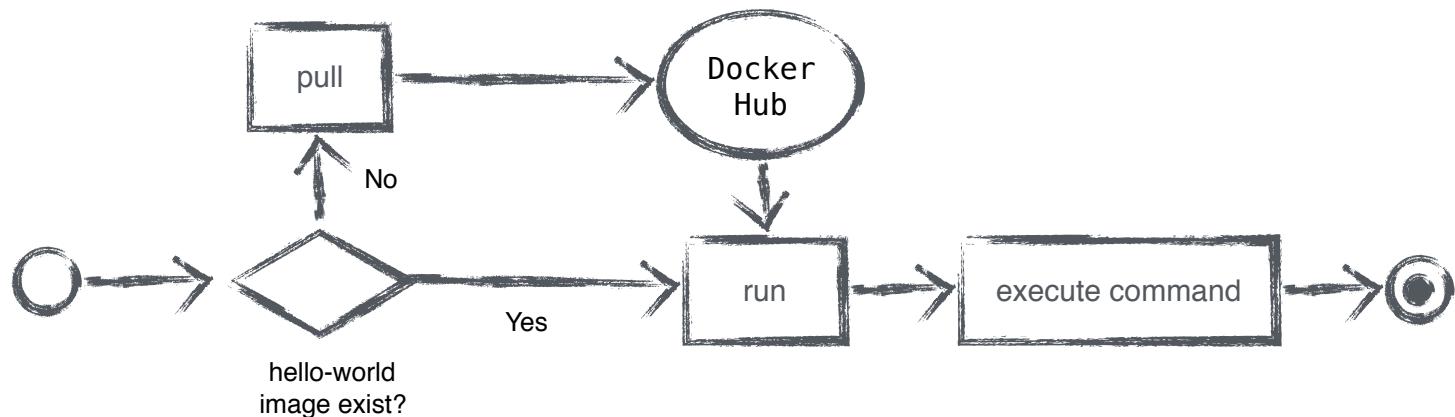
- * See the list of Docker Images.

```
$ docker images
```

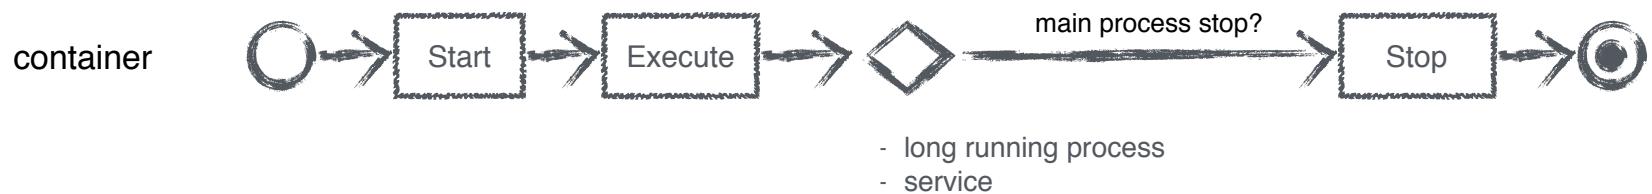
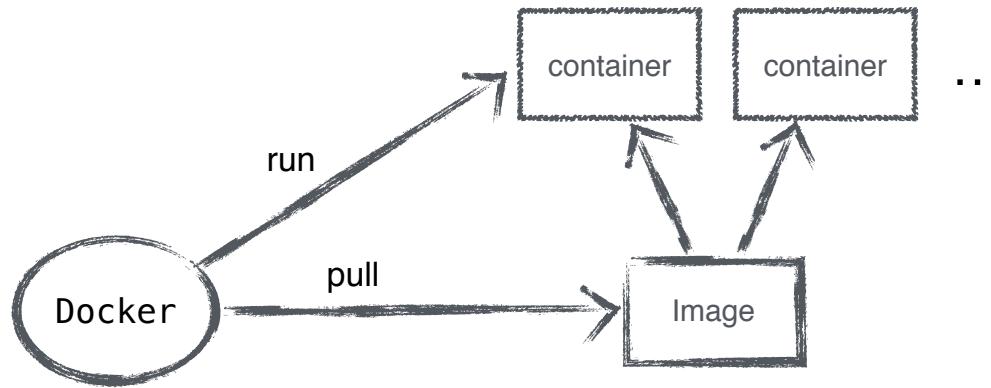
'docker run'

Usage: docker [option] COMMAND [argument ...]

```
$ docker run hello-world
```



Container State



Container management commands

command	description
<code>docker create image [command]</code> <code>docker run image [command]</code>	create the container = <code>create + start</code>
<code>docker rename container new_name</code> <code>docker update container</code>	rename the container update the container config
<code>docker start container...</code> <code>docker stop container...</code> <code>docker kill container...</code> <code>docker restart container...</code>	start the container graceful ² stop kill (SIGKILL) the container = <code>stop + start</code>
<code>docker pause container...</code> <code>docker unpause container...</code>	suspend the container resume the container
<code>docker rm [-f³] container...</code>	destroy the container

docker run — interactive mode

- * By default containers are non-interactive
 - * stdin is closed immediately
 - * terminal signals are not forwarded

```
$ docker run -t debian bash  
root@6fecc2e8ab22:/# date  
^C  
$
```

- * With -i the container runs interactively
 - * stdin is usable
 - * terminal signals are forwarded to the container

Inspecting the container

command	description
<code>docker ps</code>	list running containers
<code>docker ps -a</code>	list all containers
<code>docker logs [-f⁵] container</code>	show the container output (<i>stdout+stderr</i>)
<code>docker top container [ps options]</code>	list the processes running inside the containers ⁶
<code>docker stats [container]</code>	display live usage statistics ⁷
<code>docker diff container</code>	show the differences with the image (modified files)
<code>docker port container</code>	list port mappings
<code>docker inspect container...</code>	show low-level infos (in json format)

⁵with `-f`, `docker logs` follows the output (à la `tail -f`)

⁶`docker top` is the equivalent of the `ps` command in unix

⁷`docker stats` is the equivalent of the `top` command in unix

Interacting with the container

command	description
<code>docker attach container</code>	attach to a running container (stdin/stdout/stderr)
<code>docker cp container:path hostpath -</code> <code>docker cp hostpath - container:path</code>	copy files from the container copy files into the container
<code>docker export container</code>	export the content of the container (tar archive)
<code>docker exec container args...</code>	run a command in an existing container (useful for debugging)
<code>docker wait container</code>	wait until the container terminates and return the exit code
<code>docker commit container image</code>	commit a new docker image (snapshot of the container)

Docker Commit

```
$ docker run --name my-container -t -i debian
root@3b397d383faf:/# cat >> /etc/bash.bashrc <<EOF
> echo 'hello!'
> EOF
root@3b397d383faf:/# exit
$ docker start --attach my-container
my-container
hello!
root@3b397d383faf:/# exit
$ docker diff my-container
C /etc
C /etc/bash.bashrc
A /.bash_history
C /tmp
$ docker commit my-container hello
a57e91bc3b0f5f72641f19cab85a7f3f860a1e5e9629439007c39fd76f37c5dd
$ docker rm my-container
my-container
$ docker run --rm -t -i hello
hello!
root@386ed3934b44:/# exit
```

Docker Basics

- * In your Docker environment, just run the following command

```
$ docker run busybox echo hello world  
hello world
```

- * We used one of the smallest, simplest images available: busybox.
- * busybox is typically used in embedded systems (phones, routers...)
- * We run a single process and echo'ed hello world.

A more useful container

- * Let's run a more exciting container:

```
$ docker run -it ubuntu bash  
root@04c0bb0a6c07:/#
```

- * -it is shorthand for -i -t.
 - * -i tells Docker to connect us to the container's stdin.
 - * -t tells Docker that we want a pseudo-terminal.

- * Try to run figlet in our container:

```
root@04c0bb0a6c07:/# figlet hello  
bash: figlet: command not found  
root@04c0bb0a6c07:/# apt-get update  
root@04c0bb0a6c07:/# apt-get install figlet -y  
root@04c0bb0a6c07:/# exit  
$ docker commit 04c0bb0a6c07 hello
```

Image tags

- * A docker tag is made of two parts: “REPOSITORY:TAG”
- * The TAG part identifies the version of the image. If not provided, the default is “:latest”

```
$ docker images
REPOSITORY  TAG      IMAGE ID      CREATED      VIRTUAL SIZE
debian      8          835c4d274060  2 weeks ago  122.6 MB
debian      8.0        835c4d274060  2 weeks ago  122.6 MB
debian      jessie     835c4d274060  2 weeks ago  122.6 MB
debian      rc-buggy   350a74df81b1  7 months ago  159.9 MB
debian      experimental 36d6c9c7df4c  7 months ago  159.9 MB
debian      6.0.9      3b36e4176538  7 months ago  112.4 MB
debian      squeeze     3b36e4176538  7 months ago  112.4 MB
debian      wheezy     667250f9a437  7 months ago  115 MB
debian      latest      667250f9a437  7 months ago  115 MB
debian      7.5        667250f9a437  7 months ago  115 MB
debian      unstable    24a4621560e4  7 months ago  123.6 MB
debian      testing     7f5d8ca9fdcf  7 months ago  121.8 MB
debian      stable      caa04aa09d69  7 months ago  115 MB
debian      sid         f3d4759f77a7  7 months ago  123.6 MB
debian      7.4        e565fbcc6033  9 months ago  115 MB
debian      7.3        b5fe16f2ccba  11 months ago  117.8 MB
```

Inputs/Outputs

- * Data volumes (persistent data)
 - * mounted from the host filesystem
- * Publishing ports (NAT)

docker run — mount external volumes

- * Command

```
$ docker run -v /hostpath:/containerpath[:ro]
```

- * -v mounts the location /hostpath from the host filesystem at the location /containerpath inside the container
With the “:ro” suffix, the mount is read-only

- * Purposes:

- * store persistent data outside the container
 - * provide inputs: data, config files, ... (read-only mode)

Example:

```
docker run -i -t -v /tmp/data:/data:ro ubuntu
```

docker run — publish a TCP port

- * Command

```
$ docker run -p hostport:containerport image-name
```

- * Containers are deployed in a private network, they are not reachable from the outside (unless a redirection is set up).
- * redirect incoming connections to the TCP port hostport of the host to the TCP port containerport of the container
The listening socket binds to 0.0.0.0 (all interfaces) by default or to ipaddr if given

```
$docker run -d -p 80:80 nginx
```

Labs

1:1

- * Create nginx web server with docker on port 8080
`curl localhost:8080`

1:2

- * Install vim
- * Make changes in `/usr/share/nginx/html/index.html`
From Welcome to nginx to Welcome to KAUST

1:3

- * Commit the change to “mynginx”

1:4

- * Login to docker hub
 UserId: `kaustdocker`
 Password: `123456789`
- * Tag `mynginx` to `kaustdocker/nginx:<kaust_id>`
- * Push image to docker hub

1:5

- * Create folder `/tmp/data`
- * Mount `/tmp/data` to `/home/` in image `kaustdocker/nginx:<kaust_id>`

Dockerfile

Docker can build images automatically by reading the instructions from a **Dockerfile**. A **Dockerfile** is a text document that contains all the commands a user could call on the command line to assemble an image.

Using **docker build** users can create an automated build that executes several command-line instructions in succession.

Dockerfile

- * Dockerfile instructions are executed in order.
- * Each instruction creates a new layer in the image.
- * Instructions are cached. If no changes are detected then the instruction is skipped and the cached layer used.
- * The FROM instruction **MUST** be the first non-comment instruction.
- * Lines starting with # are treated as comments.
- * You can only have one CMD and one ENTRYPOINT instruction in a Dockerfile

Instructions affecting the image filesystem

instruction	description
<code>FROM</code> <i>image scratch</i>	base image for the build
<code>COPY</code> <i>path dst</i>	copy <i>path</i> from the context into the container at location <i>dst</i>
<code>ADD</code> <i>src dst</i>	same as <code>COPY</code> but untar archives and accepts http urls
<code>RUN</code> <i>command</i>	run an arbitrary command inside the container

Instructions setting the default container config

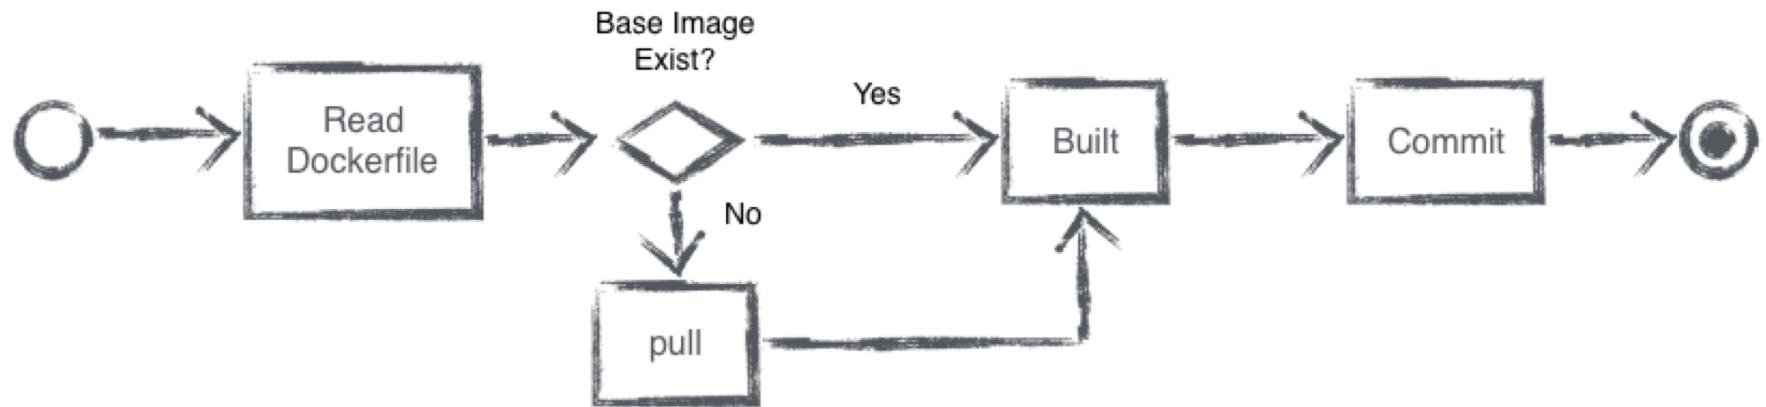
instruction	description
<code>CMD command</code> <code>ENTRYPOINT command</code>	command run inside the container entrypoint ¹³
<code>USER name[:group]</code> <code>WORKDIR path</code> <code>ENV name="value"...</code>	user running the command working directory environment variables
<code>STOP SIGNAL signal</code>	signal to be sent to terminate the container(<i>instead of SIGTERM</i>)
<code>HEALTHCHECK CMD command</code>	test command to check if the container works well
<code>EXPOSE port...</code> <code>VOLUME path...</code>	listened TCP/UDP ports mount-point for external volumes
<code>LABEL name="value"...</code>	arbitrary metadata

Dockerfile Examples

```
FROM ubuntu:14.04
MAINTAINER Hassan
RUN apt-get update && apt-get install -y nodejs npm
ADD ./src/ /src
RUN cd /src; npm install
CMD ["node", "/src/index.js"]
```

Dockerfile Usage

```
$ docker built -t my_Image .
```



Dockerfile

Multi-stage build

```
#===== Stage 1: build the app from sources =====#
FROM debian:stretch AS builder
# update the package lists and install the build dependencies
RUN apt-get -qqy update
RUN apt-get -qqy install gcc make libacme-dev

# install the sources in /opt/src and build them
COPY . /opt/src
RUN cd /opt/src && ./configure && make
# install the files in a tmp dir and make an archive that we can
# deploy elsewhere
RUN cd /opt/src && make install DESTDIR=/tmp/dst \
&& cd /tmp/dst && tar czvf /tmp/myapp.tgz .

#===== Stage 2: final image =====#
FROM debian:stretch
# update the package lists and install the runtime dependencies
RUN apt-get -qqy update
RUN apt-get -qqy install libacme1.0
# install the app built in stage 1
COPY --from=builder /tmp/myapp.tgz /tmp/
RUN cd / && tar zxf /tmp/myapp.tgz
CMD ["myapp"]
```

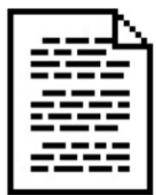
Labs

2:1

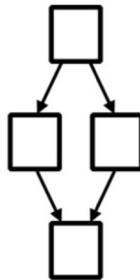
- * Using Dockerfile
 - * build nginx web server with docker on port 8080 and tag it “mynginx:lab2”
 - * Install vim
 - * Copy
 - <https://raw.githubusercontent.com/kaustdocker/docker101/master/index.html>
 - * to /usr/share/nginx/html/

Docker Compose

Compose is a tool for defining and running multi-container Docker applications. With Compose, you use a YAML file to configure your application's services. Then, with a single command, you create and start all the services from your configuration



`docker-compose up`



Text file

group.yml

```
name: counter

containers:
  web:
    build: .
    command: python app.py
    ports:
      - "5000:5000"
    volumes:
      - .:/code
    links:
      - redis
  redis:
    image: redis:latest
```

Docker Compose

Commands

```
docker-compose start
docker-compose stop
docker-compose pause
docker-compose unpause
docker-compose ps
docker-compose up
docker-compose down
```

Docker Compose

Reference

Building

```
web:  
  # build from Dockerfile  
  build: .  
  
  # build from custom Dockerfile  
  build:  
    context: ./dir  
    dockerfile: Dockerfile.dev  
  
  # build from image  
  image: ubuntu  
  image: ubuntu:14.04  
  image: tutum/influxdb
```

Commands

```
# command to execute  
command: bundle exec thin -p 3000  
  command: [bundle, exec, thin, -p, 3000]  
  
# override the entrypoint  
entrypoint: /app/start.sh  
  entrypoint: [php, -d, vendor/bin/phpunit]
```

Ports

```
ports:  
  - "3000"  
  - "8000:80" # guest:host  
# expose ports to linked services (not to host)  
expose: ["3000"]
```

Environment variables

```
# environment vars  
environment:  
  RACK_ENV: development  
environment:  
  - RACK_ENV=development  
  
# environment vars from file  
env_file: .env  
env_file: [.env, .development.env]
```

Docker Compose Examples

```
version: '3.3'

services:
  db:
    image: mysql:5.7
    volumes:
      - db_data:/var/lib/mysql
    restart: always
    environment:
      MYSQL_ROOT_PASSWORD: somewordpress
      MYSQL_DATABASE: wordpress
      MYSQL_USER: wordpress
      MYSQL_PASSWORD: wordpress

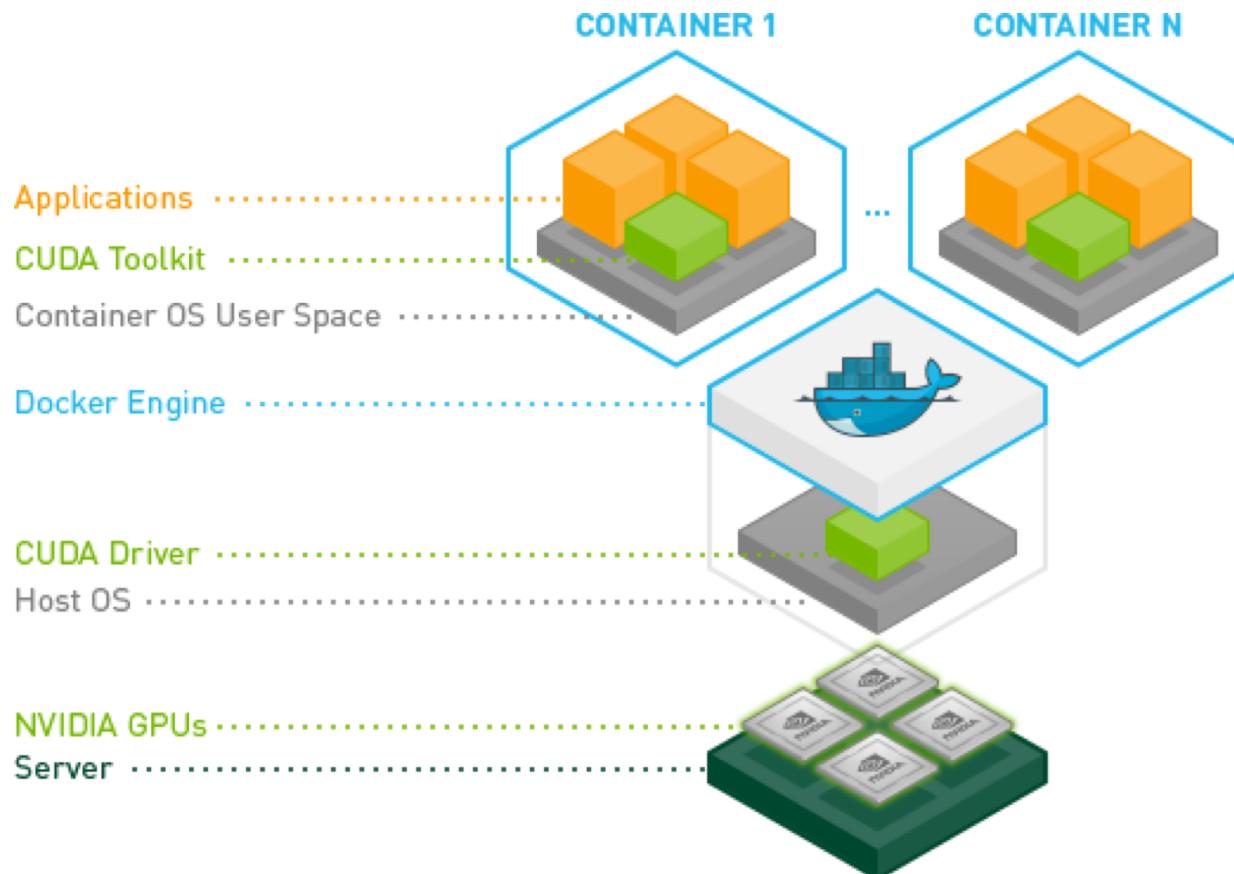
  wordpress:
    depends_on:
      - db
    image: wordpress:latest
    ports:
      - "8000:80"
    restart: always
    environment:
      WORDPRESS_DB_HOST: db:3306
      WORDPRESS_DB_USER: wordpress
      WORDPRESS_DB_PASSWORD: wordpress
      WORDPRESS_DB_NAME: wordpress
  volumes:
    db_data: {}
```

Exercise

<https://docs.docker.com/compose/gettingstarted/>

<https://github.com/kaustdocker/docker101/tree/master/composetest>

NVIDIA Container Runtime for Docker



NVIDIA Container Runtime for Docker Installation

```
# Add the package repositories
curl -s -L https://nvidia.github.io/nvidia-docker/gpgkey | \
  sudo apt-key add -
distribution=$(./etc/os-release;echo $ID$VERSION_ID)
curl -s -L https://nvidia.github.io/nvidia-docker/$distribution/nvidia-docker.list | \
  sudo tee /etc/apt/sources.list.d/nvidia-docker.list
sudo apt-get update

# Install nvidia-docker2 and reload the Docker daemon configuration
sudo apt-get install -y nvidia-docker2
sudo pkill -SIGHUP dockerd
```

NVIDIA Container Runtime for Docker

Usage

In docker command :

```
docker run --runtime=nvidia --rm nvidia/cuda:9.0-base nvidia-smi
```

In docker-compose :

```
---
```

```
version: "2.3"
services:
  paraview:
    hostname: paraview
    tty: true
    image: kcr.kaust.edu.sa/apps/paraview
    runtime: nvidia
    environment:
      - DISPLAY=unix${DISPLAY}
      - NVIDIA_VISIBLE_DEVICES=all
    working_dir:
      "/home/${USER}/"
    volumes:
      - "/var/lib/sss/pipes/:/var/lib/sss/pipes/:rw"
      - "/home/${USER}/:/home/${USER}/"
      - "/tmp/.X11-unix:/tmp/.X11-unix:rw"
    command:
      /bin/su - $USER -c "/opt/ParaView-5.6.0-MPI-Linux-64bit/bin/paraview"
```

Container orchestration

Container orchestration is all about managing the lifecycles of containers, especially in large, dynamic environments. Software teams use container orchestration to control and automate many tasks:

- * Provisioning and deployment of containers
- * Redundancy and availability of containers
- * Scaling up or removing containers to spread application load evenly across host infrastructure
- * Movement of containers from one host to another if there is a shortage of resources in a host, or if a host dies
- * Allocation of resources between containers
- * External exposure of services running in a container with the outside world
- * Load balancing of service discovery between containers
- * Health monitoring of containers and hosts
- * Configuration of an application in relation to the containers running it

CONTAINER PLATFORM COMPARISON (06/2017)

	Kubernetes	OpenShift	DC/OS	ECS	Docker	Nomad	Cattle	Kontena
Memory	✓	✓	✓	✓	✓	✓	✓	✓
CPU	✓	✓	✓	✓	✓	✓	✓	✓
GPU	⊖	⊖	⊖					
Disk Space	⊖		✓					
Ephemeral Volumes	✓	✓	✓	✗	✓	⊖		✓
Remote Persistent Vols	⊖	⊖	⊖					
Local Persistent Vols	⊖		✓					
Ports	✓	✓	✓	✓	✓	✓	✓	✓
IPs (per container)	⊖	✓	⊖		⊖		⊖	⊖

 Included  New/External/Partial/Experimental

CONTAINER PLATFORM COMPARISON (06/2017)

	Kubernetes	OpenShift	DC/OS	ECS	Docker	Nomad	Cattle	Kontena
Placement	✓	✓	✓	✓	✓	✓	✓	✓
Replication/Scaling	✓	✓	✓	✓	✓	✓	✗	✓
Readiness Checking	✓	✓	✓	✗	✗	✓	✓	✓
Resurrection	✓	✓	✗	✗	✓	✓	✓	✓
Rescheduling	✓	✓	✓	✓	✓	✗	✓	✓
Rolling Updates	✓	✓	✓	✗	✗	✓	✓	✗
Collocation	✓	✓	✗			✗	✗	
Daemons	✗	✗		✗				
Cron Jobs	✗	✗	✗	✗		✗		

✓ Included

✗ New/External/Partial/Experimental

Nvidia, TensorFlow & Kubernetes